End-to-End Protocols: UDP and TCP

Hui Chen, Ph.D.

Dept. of Engineering & Computer Science

Virginia State University

Petersburg, VA 23806

Acknowledgements

- Some pictures used in this presentation were obtained from the Internet
- The instructor used the following references
 - Larry L. Peterson and Bruce S. Davie, Computer Networks: A Systems Approach, 5th Edition, Elsevier, 2011
 - Andrew S. Tanenbaum, Computer Networks, 5th Edition, Prentice-Hall, 2010
 - James F. Kurose and Keith W. Ross, Computer Networking: A Top-Down Approach, 5th Ed., Addison Wesley, 2009
 - Larry L. Peterson's (http://www.cs.princeton.edu/~llp/) Computer
 Networks class web site

Acknowledgements

- Animations in the PDF version of the slides is produced using
 - PPspliT
 - http://www.dia.uniroma3.it/~rimondin/downloads.php

Outline

- User Datagram Protocol
- **□** Transmission Control Protocol

Network Applications



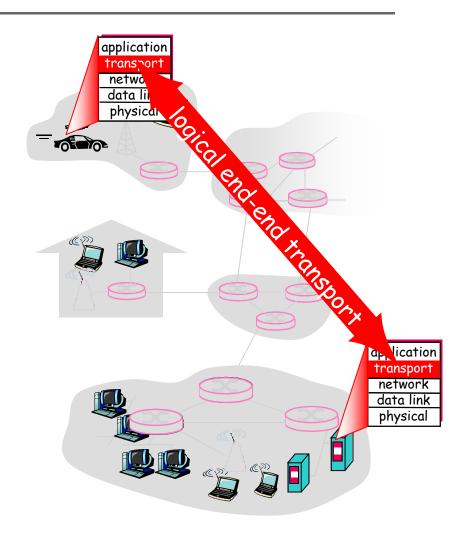


- Users make use of networks via network applications at hosts
- A hosts can run many network applications simultaneously
- Each application is one or more running programs (processes)
- Q: How processes share the underlying network layers?



Transport Layer Services and Protocols

- provide logical communication between application processes running on different hosts
- transport protocols run in end systems
 - send side
 - breaks app messages into segments, passes to network layer
 - receive side:
 - reassembles segments into messages, passes to app layer
- more than one transport protocol available to applications
 - Internet: TCP and UDP



Transport vs. Network Layer (1)

- network layer: logical communication between hosts
- transport layer: logical communication between processes
 - relies on, enhances, network layer services

Household analogy:

- 12 kids sending letters among themselves via their parents
- processes = kids
- application messages = letters in envelopes
- □ hosts = houses
- transport protocol = Ann and Bill (parents)
- network-layer protocol = postal service

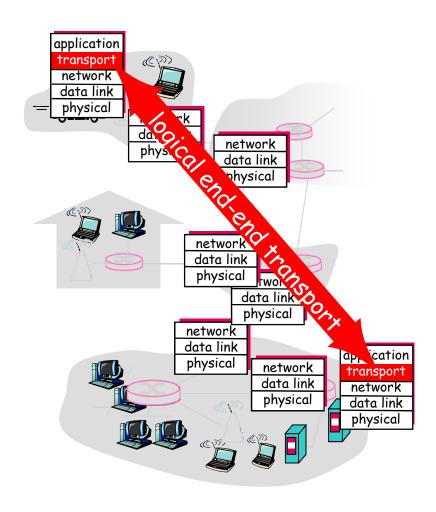
Transport vs. Network Layer (2)

- Network layer: Underlying besteffort network
 - drop messages
 - re-orders messages
 - delivers duplicate copies of a given message
 - limits messages to some finite size
 - delivers messages after an arbitrarily long delay

- Transport Layer: Common endto-end services
 - guarantee message delivery
 - deliver messages in the same order they are sent
 - deliver at most one copy of each message
 - support arbitrarily large messages
 - support synchronization
 - allow the receiver to flow control the sender
 - support multiple application processes on each host

Internet Transport-Layer Protocols

- Reliable, in-order delivery (TCP)
 - congestion control
 - flow control
 - connection setup
- Unreliable, unordered delivery: UDP
 - no-frills extension of "best-effort" IP
- Services not available:
 - delay guarantees
 - bandwidth guarantees



Multiplexing/Demultiplexing

Host-to-host delivery ←→ process-to-process delivery

Multiplexing/Demultiplexing

Host-to-host delivery ←→ process-to-process delivery

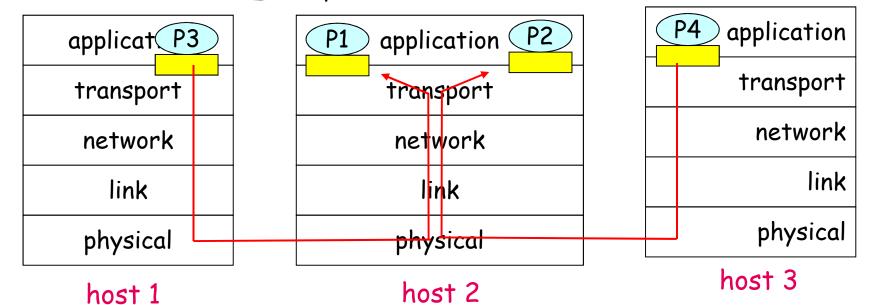
Demultiplexing at rcv host:

delivering received segments to correct socket

= socket = process

Multiplexing at send host:


gathering data from multiple sockets, enveloping data with header (later used for demultiplexing)



Simple Demultiplexer (1)

- Need to know to or from which process the data is sent or come
 - Identify processes on hosts
- How to identify processes on hosts?
 - Introduce concept of "port"
 - Q: why not to use process id?

Processes ID: Windows Example




Processes ID: Linux Example

```
🧬 hchen@turing:~
[hchen@turing ~] $ ps ax
  PID TTY
               STAT
                       TIME COMMAND
    1 ?
               Ss
                       0:02 init [5]
                      0:00 [migration/0]
               S<
               sn
                      0:00 [ksoftirqd/0]
                      0:00 [watchdog/0]
               ≲<
                      0:00 [migration/1]
               ಽ<
                      0:00 [ksoftirqd/1]
               sn
                      0:00 [watchdog/1]
               S<
    8 ?
               S<
                      0:00 [migration/2]
    9 ?
                      0:00 [ksoftirqd/2]
               sn
                      0:00 [watchdog/2]
   10 ?
               8<
                      0:00 [migration/3]
   11 ?
               8<
   12 ?
                      0:00 [ksoftirqd/3]
               sn
   13 ?
               8<
                      0:00 [watchdog/3]
                      0:00 [migration/4]
   14 ?
               8<
                      0:00 [ksoftirqd/4]
   15 ?
               sn
   16 ?
               8<
                      0:00 [watchdog/4]
                      0:00 [migration/5]
   17 2
               ಽ<
   18 ?
               sn
                      0:00 [ksoftirqd/5]
                      0:00 [watchdog/5]
   19 ?
               s≺
                      0:00 [migration/6]
   20 2
               s≺
                       0:00 [ksoftirqd/6]
   21 2
               sn
                       0:00 [watchdog/6]
               8<
   22 ?
```

Simple Demultiplexer (2)

- How to identify processes on hosts?
 - Q: why not to use process id?
 - Introduce concept of "port"
 - Endpoints identified by ports
 - servers have well-known ports
 - see /etc/services on Unix/Linux
 - □ see C:\WINDOWS\system32\\drivers\etc\services on MS Windows




10/31/2016

CSCI 445 - Fall 2016

Simple Demultiplexer: UDP


- Adds multiplexing to Internet Protocol
 - Endpoints identified by ports (UDP ports)
 - Demultiplex via ports on hosts
 - Nothing more is added
 - Unreliable and unordered datagram service
 - No flow control
 - User Datagram Protocol (UDP)
 - A process is identified by <host, port>
 - Connectionless model
- Header format
 - Optional checksum
 - psuedo header + UDP header + data
 - pseudo header = protocol number + source IP address and destination IP address + UDP length field



From UDP header

From IP header

In-Class Exercise L15-1



- Q1: How many UDP ports are there?
- Q2: How big are UDP headers?
- Q3: How much data does a UDP datagram can carry?
- Turn your work in before you leave!


Transmission Control Protocol (TCP)

- Connection-oriented
- Byte-stream
 - applications writes bytes
 - TCP sends segments
 - applications reads bytes
- □ Full duplex
- ☐ Flow control: keep sender from overrunning receiver
- Congestion control: keep sender from overrunning network


Data Link Versus Transport

- Potentially connects many different hosts
 - need explicit connection establishment and termination
- Potentially different RTT
 - need adaptive timeout mechanism
- Potentially long delay in network
 - need to be prepared for arrival of very old packets

- Potentially different capacity at destination
- need to accommodate different node capacity
- Potentially different network capacity
- need to be prepared for network congestion




Segment Format (1)



Segment Format (2)

- Each connection identified with 4-tuple:
 - (SrcPort, SrcIPAddr, DsrPort, DstIPAddr)
- Sliding window + flow control
 - acknowledgment, SequenceNum, AdvertisedWinow
- Flags
 - SYN, FIN, RESET, PUSH, URG, ACK
- Checksum
 - pseudo header + TCP header + data



Sequence and Acknowledgement Numbers (1)

- Host A sends a file of 500,000 bytes over a TCP connection with Maximum Segment Size (MSS) as 1,000 bytes to host B
 - How many segments? 500,000/1,000 = 500
 - Sequence number assignments
 - Sequence number of 1st segment? 0
 - □ Sequence number of 2nd segment? 1,000
 - □ Sequence number of 3rd segment? 2,000
 - **-**

Sequence and Acknowledgement Numbers (2)

■ Scenario 1

- Host B received all bytes numbered 0 to 1,999 from host A
- What would host B put in the acknowledgement number field of the segment it sends to A?
 - 2,000: the sequence number of the next byte host B is expecting

Scenario 2

- Host B received two segments containing bytes from 0-999, and 2,000-2,999, respectively?
- What would host B put in the acknowledgement number field of the segment it sends to A?
 - 1000: TCP only acknowledges bytes up to the first missing byte in the stream, and it is the next byte host B is expecting


■ Scenario 3

- Host B received 1st segment containing bytes from 0-999. Somehow, next it received 3rd segment containing bytes from 2,000-2,999.
- What does host B in this case that the segments arrive out of order?
 - □ TCP does not specify how to deal with this situation. Hence, it is up to the implementation.
 - Option 1: Host B immediately discards out-of-order segment → simple receiver design
 - Option 2: Host B keeps the out-of-order segment and waits for missing bytes to fill in the gaps → more efficient on bandwidth utilization → taken in practice


TCP is Connection-Oriented

- Keep track of states of receiver and sender
 - Connection Establishment
 - Connection Termination
 - TCP finite state machine and state transition

Connection Establishment




Connection Termination




State Transition Diagram


10/31/2016





10/31/2016 CSCI 445 – Fall 2016 26


















































































































Connection Termination and State Transition (2)

- ☐ This side closes first
 - ESTABLISHED → FIN_WAIT_1 → FIN_WAIT_2 → TIME_WAIT
- Other side closes first
 - ESTABLISHED → CLOSE_WAIT → LAST_ACK → CLOSED
- Both sides close at the same time
 - ESTABLISHED → FIN_WAIT_1 → CLOSING → TIME_WAIT
 → CLOSED

TCP Sliding Window: Why Different?


- Potentially connects many different hosts
 - need explicit connection establishment and termination
- Potentially different RTT
 - need adaptive timeout mechanism
- Potentially long delay in network
 - need to be prepared for arrival of very old packets


- Potentially different capacity at destination
 - need to accommodate different node capacity
- Potentially different network capacity
 - need to be prepared for network congestion



TCP Sliding Window: Reliable and Ordered Delivery

TCP uses cumulative acknowledgements to acknowledge receiving of all the bytes up to the first missing byte






- Sending side
 - LastByteAcked ≤ LastByteSent
 - LastByteSent ≤ LastByteWritten
 - buffer bytes between LastByteAcked and LastByteWritten

Receiving side
LastByteRead < NextByteExpected
NextByteExpected ≤ LastByteRcvd +1
buffer bytes betweenNextByteRead and
LastByteRcvd

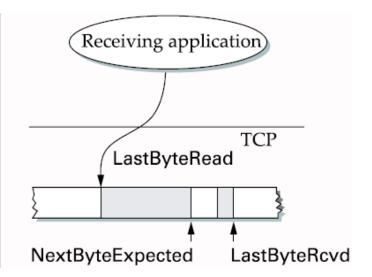
TCP Flow Control (1)

- receive side of TCP connection has a receive buffer
- app process may be slow at reading from buffer
- speed-matching service: matching the send rate to the receiving app's drain rate



flow control-

sender won't overflow receiver's buffer by transmitting too much, too fast



TCP Flow Control (2)

- Send buffer size: MaxSendBuffer
- Receive buffer size: MaxRcvBuffer
- Receiving side
 - LastByteRcvd LastByteRead ≤ MaxRcvBuffer
 - AdvertisedWindow = MaxRcvBuffer ((NextByteExpected -1) -LastByteRead)) → maximum possible free space remaining in the buffer
- Sending side
 - LastByteSent LastByteAcked ≤ AdvertisedWindow
 - LastByteSent LastByteAcked: unacknowledged bytes sender has put in TCP
 - Otherwise, the sender may overrun the receiver
 - EffectiveWindow = AdvertisedWindow (LastByteSent -LastByteAcked)
 → how much data it can sent
 - LastByteWritten LastByteAcked ≤ MaxSendBuffer
 - If the sender tries to write y bytes to TCP
 - □ block sender if (LastByteWritten LastByteAcked) + y > MaxSenderBuffer
- Always send ACK in response to arriving data segment
- Persist when AdvertisedWindow = 0

Flow Control and Buffering (3)

_	<u>A</u>	Message	B	Comments
1	-	< request 8 buffers>	-	A wants 8 buffers
2	•	<ack 15,="" =="" buf="4"></ack>	•	B grants messages 0-3 only
3		<seq 0,="" =="" data="m0"></seq>	-	A has 3 buffers left now
4		<seq 1,="" =="" data="m1"></seq>	-	A has 2 buffers left now
5		<seq 2,="" =="" data="m2"></seq>	•••	Message lost but A thinks it has 1 left
6	•	<ack = 1, buf = 3>	•	B acknowledges 0 and 1, permits 2-4
7		<seq = 3, data = m3 $>$		A has 1 buffer left
8		<seq 4,="" =="" data="m4"></seq>		A has 0 buffers left, and must stop
9		<seq = 2, data = m2>		A times out and retransmits
10	•	<ack = 4, buf = 0>	←	Everything acknowledged, but A still blocked
11	•	<ack = 4, buf = 1>	←	A may now send 5
12	•	<ack = 4, buf = 2 $>$	•	B found a new buffer somewhere
13		<seq = 5, data = m5 $>$		A has 1 buffer left
14		<seq = 6, data = m6 $>$	-	A is now blocked again
15	•	<ack = 6, buf = 0>	•	A is still blocked
16	•••	<ack = 6, buf = 4 $>$	•	Potential deadlock
D,	mamic h	uffor allocation. The arrow	c chow t	ha direction of transmission. An allinsis /

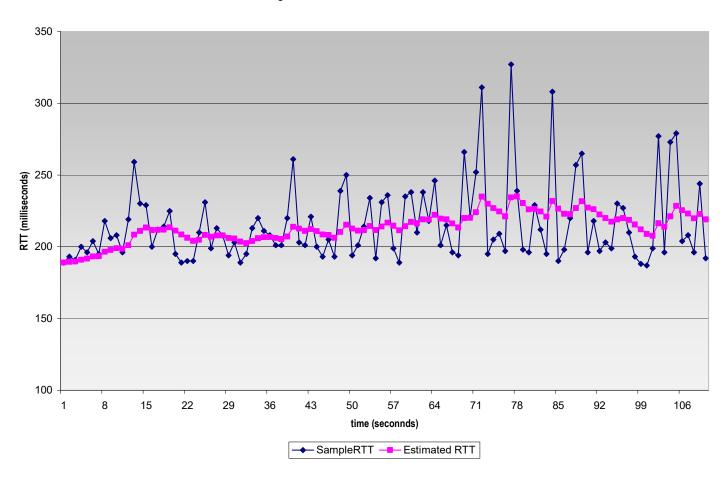
Dynamic buffer allocation. The arrows show the direction of transmission. An ellipsis (...) indicates a lost TCP segment

Adaptive Retransmission: Original Algorithm

- Measure SampleRTT for each segment/ACK pair
- □ Compute weighted average of RTT
 - **E**stimatedRTT = α x EstimatedRTT + β x SampleRTT
 - where $\alpha + \beta = 1$
 - \square α between 0.8 and 0.9
 - β between 0.1 and 0.2
 - Set timeout based on EstimatedRTT
 - □ TimeOut = 2 x EstimatedRTT

Example RTT estimation:

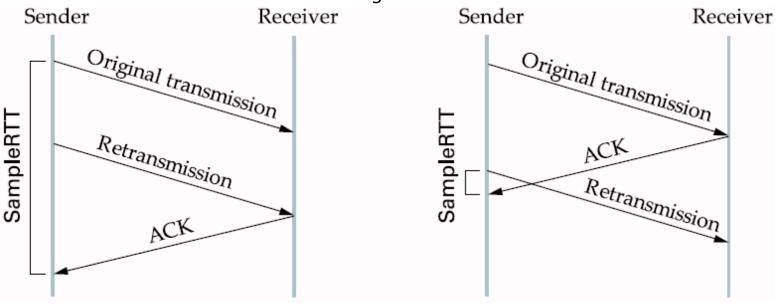
RTT: gaia.cs.umass.edu to fantasia.eurecom.fr



Adaptive Retransmission: Karn/Partridge Algorithm

Problem with original algorithm

ACK does not really acknowledge a transmission, it acknowledges the receipt of data → can not distinguish an ACK is for which transmission/retransmission of a segment



- Do not sample RTT when retransmitting
- Double timeout after each retransmission
 - Congestion is the most likely cause of lost segments → TCP should not react too aggressively to a timeout

Jacobson/ Karels Algorithm

- Previous approaches did not take the variance of the sample RTT into account
 - If no variance, Estimated RTT is good enough, $2 \times Estimated RTT$ is too pessimistic
 - If variance large, timeout value should not be too dependent on Estimated RTT
- New Calculations for average RTT
 - Difference = SampleRTT EstimtaedRTT
 - EstimatedRTT = EstimatedRTT + (δ x Difference)
 - Deviation = Deviation + δ (| Difference | Deviation)
 - \square where δ is a factor between 0 and 1
 - Consider variance when setting timeout value
 - □ TimeOut = μ x EstimatedRTT + ϕ x Deviation
 - \square where $\mu = 1$ and $\varphi = 4$
- Notes
 - algorithm only as good as granularity of clock (500ms on Unix)
 - accurate timeout mechanism important to congestion control

TCP: Sequence Number Wrap Around

Bandwidth	Time until Wraparound
T1 (1.5 Mbps)	6.4 hours
Ethernet (10 Mbps)	57 minutes
T3 (45 Mbps)	13 minutes
Fast Ethernet (100 Mbps)	6 minutes
OC-3 (155 Mbps)	4 minutes
OC-12 (622 Mbps)	55 seconds
OC-48 (2.5 Gbps)	14 seconds

Time until 32-bit sequence number space wraps around

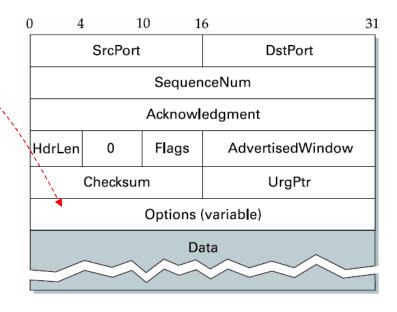
TCP: Can Keep Pipe Full?

Bandwidth	Delay × Bandwidth Product	
T1 (1.5 Mbps)	18 KB	
Ethernet (10 Mbps)	122 KB	
T3 (45 Mbps)	549 KB	
Fast Ethernet (100 Mbps)	1.2 MB	
OC-3 (155 Mbps)	1.8 MB	
OC-12 (622 Mbps)	7.4 MB	
OC-48 (2.5 Gbps)	29.6 MB	

Required window size for 100-ms RTT.

Solution: TCP Extensions

- Implemented as header options
- Store timestamp in outgoing segments → measure RTT
- Extend sequence space with 32bit timestamp → protected against sequence number wraparound
- □ Shift (scale) advertised window → keep the pipe full
- □ Selective acknowledgement (SAC)
 → acknowledge any additional (out-of-order) blocks of received data



TCP Extensions for High Performance http://tools.ietf.org/html/rfc1323

Summary

- User Datagram Protocol
 - Multiplexer/Demultiplexer for IP
- Transmission Control Protocol
 - Reliable Byte Stream
 - Connection-oriented
 - Connection establishment
 - Connection termination
 - Automatics Repeated-Request: ACKs and NACKs
 - □ Flow-control
 - □ Timeout value estimation
 - Extensions
- Congestion control (future discussions)